

Statistical Machine Learning

Exercise sheet 6

Exercise 6.1 (Multiclass Logistic Regression) In this exercise we derive a multiclass generalization of logistic regression. We shall assume that the input variable X is a vector in \mathbb{R}^p as before, but the output variable Y is a vector of the form $y = (y_1, \dots, y_K) \in \{0, 1\}^K$ with $\sum_{k=1}^K y_k = 1$. Thus, the vector y such that $y_k = 1$ for $k = m$ and 0 otherwise, corresponds to the class m . It follows that the probability of X being in class m given $X = x$ is $\mathbb{P}(Y_m = 1|X = x)$, where $Y = (Y_1, \dots, Y_K)$.

- (a) Let $w_1, \dots, w_K \in \mathbb{R}^p$ be K vectors of parameters each associated with the corresponding class. Construct a conditional model for $Y = y|X = x$ such that $\mathbb{P}(Y_k = 1|X = x) \propto \exp(w_k^\top x)$. In particular, find $\mathbb{P}(Y_k = 1|X = x)$.
- (b) Show that when $K = 2$, the proposed model is equivalent to logistic regression, except that the model is over-parameterized, and therefore w_1 and w_2 are not identifiable. Is this a problem?
- (c) Show that the model is still overparametrized if $K > 2$ and that one can impose the constraint $\sum_{k=1}^K w_k = 0$.
- (d) Express $\mathbb{P}(Y_k = 1|Y_k + Y_j = 1, X = x)$, or alternatively, derive the log-odds between two classes. What is the shape of $\{x \mid \mathbb{P}(Y_k = 1|X = x) = \mathbb{P}(Y_j = 1|X = x)\}$? Deduce that the region of space where class k is most likely is a polyhedron.
- (e) Assume that we have a sample $\{(x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)})\}$ with $(x^{(i)} \in \mathbb{R}^p$ and $y^{(i)}$ an indicator vector. Write the negative (conditional) log-likelihood of the sample and show that it can be interpreted as the empirical risk associated with a loss function that you will specify $\ell : \mathbb{R}^K \times \{0, 1\}^K \rightarrow \mathbb{R}$ applied to a predictor $f(x)$ of the form $f(x) = (f_1(x), \dots, f_K(x))$ with $f_k(x) = w_k^\top x$.
- (f) How would you apply Tikhonov regularization to the corresponding empirical risk?
- (g) Since the model is over-parameterized, instead of using the strategy proposed in (c), one could propose to just set $w_K = 0$. Prove that this would yield an equivalent model.
- (h) Now, if we use Tikhonov regularization, why is the option to set $w_K = 0$ not such a good idea?
- (i) Why is the option proposed in (c) better? If we regularize with Tikhonov regularization and don't enforce the constraint $\sum_{k=1}^K w_k = 0$, what happens?

Practical exercises

Exercise 6.2 (Implementation of the LDA and QDA algorithms and comparison with logistic regression) The files `classificationA.train`, `classificationB.train` and `classificationC.train` contain samples of data (x_i, y_i) where $x_i \in \mathbb{R}^2$ and $y_i \in \{0, 1\}$ (each line of each file contains the 2 components of x_i then y_i). The goal of this exercise is to implement linear classification methods and to test them on the three data sets.

- (a) For each data set (A,B,C) represent graphically the training data as a point cloud in \mathbb{R}^2 using different markers for the two classes using `ggplot`,

```
#Input
inp <- scan("classificationA.train", list(x1=0,x2=0,y=0))
inp <- data.frame(x1=inp$x1,x2=inp$x2,y=inp$y)
#Base Plot
G <- ggplot(data = inp, mapping = aes(x = x1, y = x2,
  color = as.factor(y)))+ geom_point()
```

- (b) Apply LDA and compute the MLE estimates for all the parameters. Plot the classification boundary for LDA for each data set by completing the following R code by entering correct values for `b` and `S` which are determined by the fact that the boundary is given by $w[1]x_1 + w[2]x_2 + b = 0$.

```
#LDA Boundary Parameters
b <- << ENTER b HERE >>
w <- << ENTER w HERE >>

#LDA Boundary Function:  $x_2 = (-x_1 \cdot w[1] + b) / w[2]$ 
LDAcurve <- function(x) {
  (- x*w[1] + b)/w[2]
}
#Plot with LDA boundary
Graph_LDA <- G
+ stat_function(fun = LDAcurve, color = "black")
```

- (c) On a separate figure, plot the classification boundary for QDA, on top of the data for each data set by completing the following code. Because we do not have a handy expression for the graph of the boundary, say $f(x_1, x_2) = 0$, we shall draw it as a contour of $f(x_1, x_2) = z$. Enter the expression from `f(x1,x2)` below.

```
#QDA Contour
cont_QDA <- curve3d(<< ENTER FUNCTION f(x1,x2) HERE >>,
  from = c(-6,-6), to = c(6,6), n=c(100,100),
  sys3d="none")
dimnames(cont_QDA$z) <- list(cont_QDA$x, cont_QDA$y)
M_QDA <- reshape2::melt(cont_QDA$z)
#Plot with QDA boundary
Graph_QDA <- G
```

```
+ geom_contour(data=M_QDA,
aes(x=Var1,y=Var2,z=value),
breaks=0, linejoin = "round", colour="black")
```

- (d) Run logistic regression using the `glm` function in R, and make again a similar plot with the data and the decision boundary using `stat_function` in `ggplot`.

```
#Logistic Regression
logres <- glm(y ~ x1 + x2, data = inp, family = binomial)
summary(logres)$coef

#Logistic Regression Coefficients
m1 <- summary(logres)$coef[[2,1]] #Coefficient of x1
m2 <- summary(logres)$coef[[3,1]] #Coefficient of x2
mc <- summary(logres)$coef[[1,1]] #Constant term

#Logistic Regression Boundary Function:  $f(x) = -(mc + m1 * x1)/m2$ 
logcurve <- function(x) {
  << ENTER CODE HERE >>
}
```

Are the coefficients very large? If so, why?

- (e) Do the same visualizations on the three testing data sets.
- (f) Compute the misclassification error of all three methods on all the three training sets and their corresponding testing sets. Which method performs better and why?